



Node.js Microservices on Autopilot

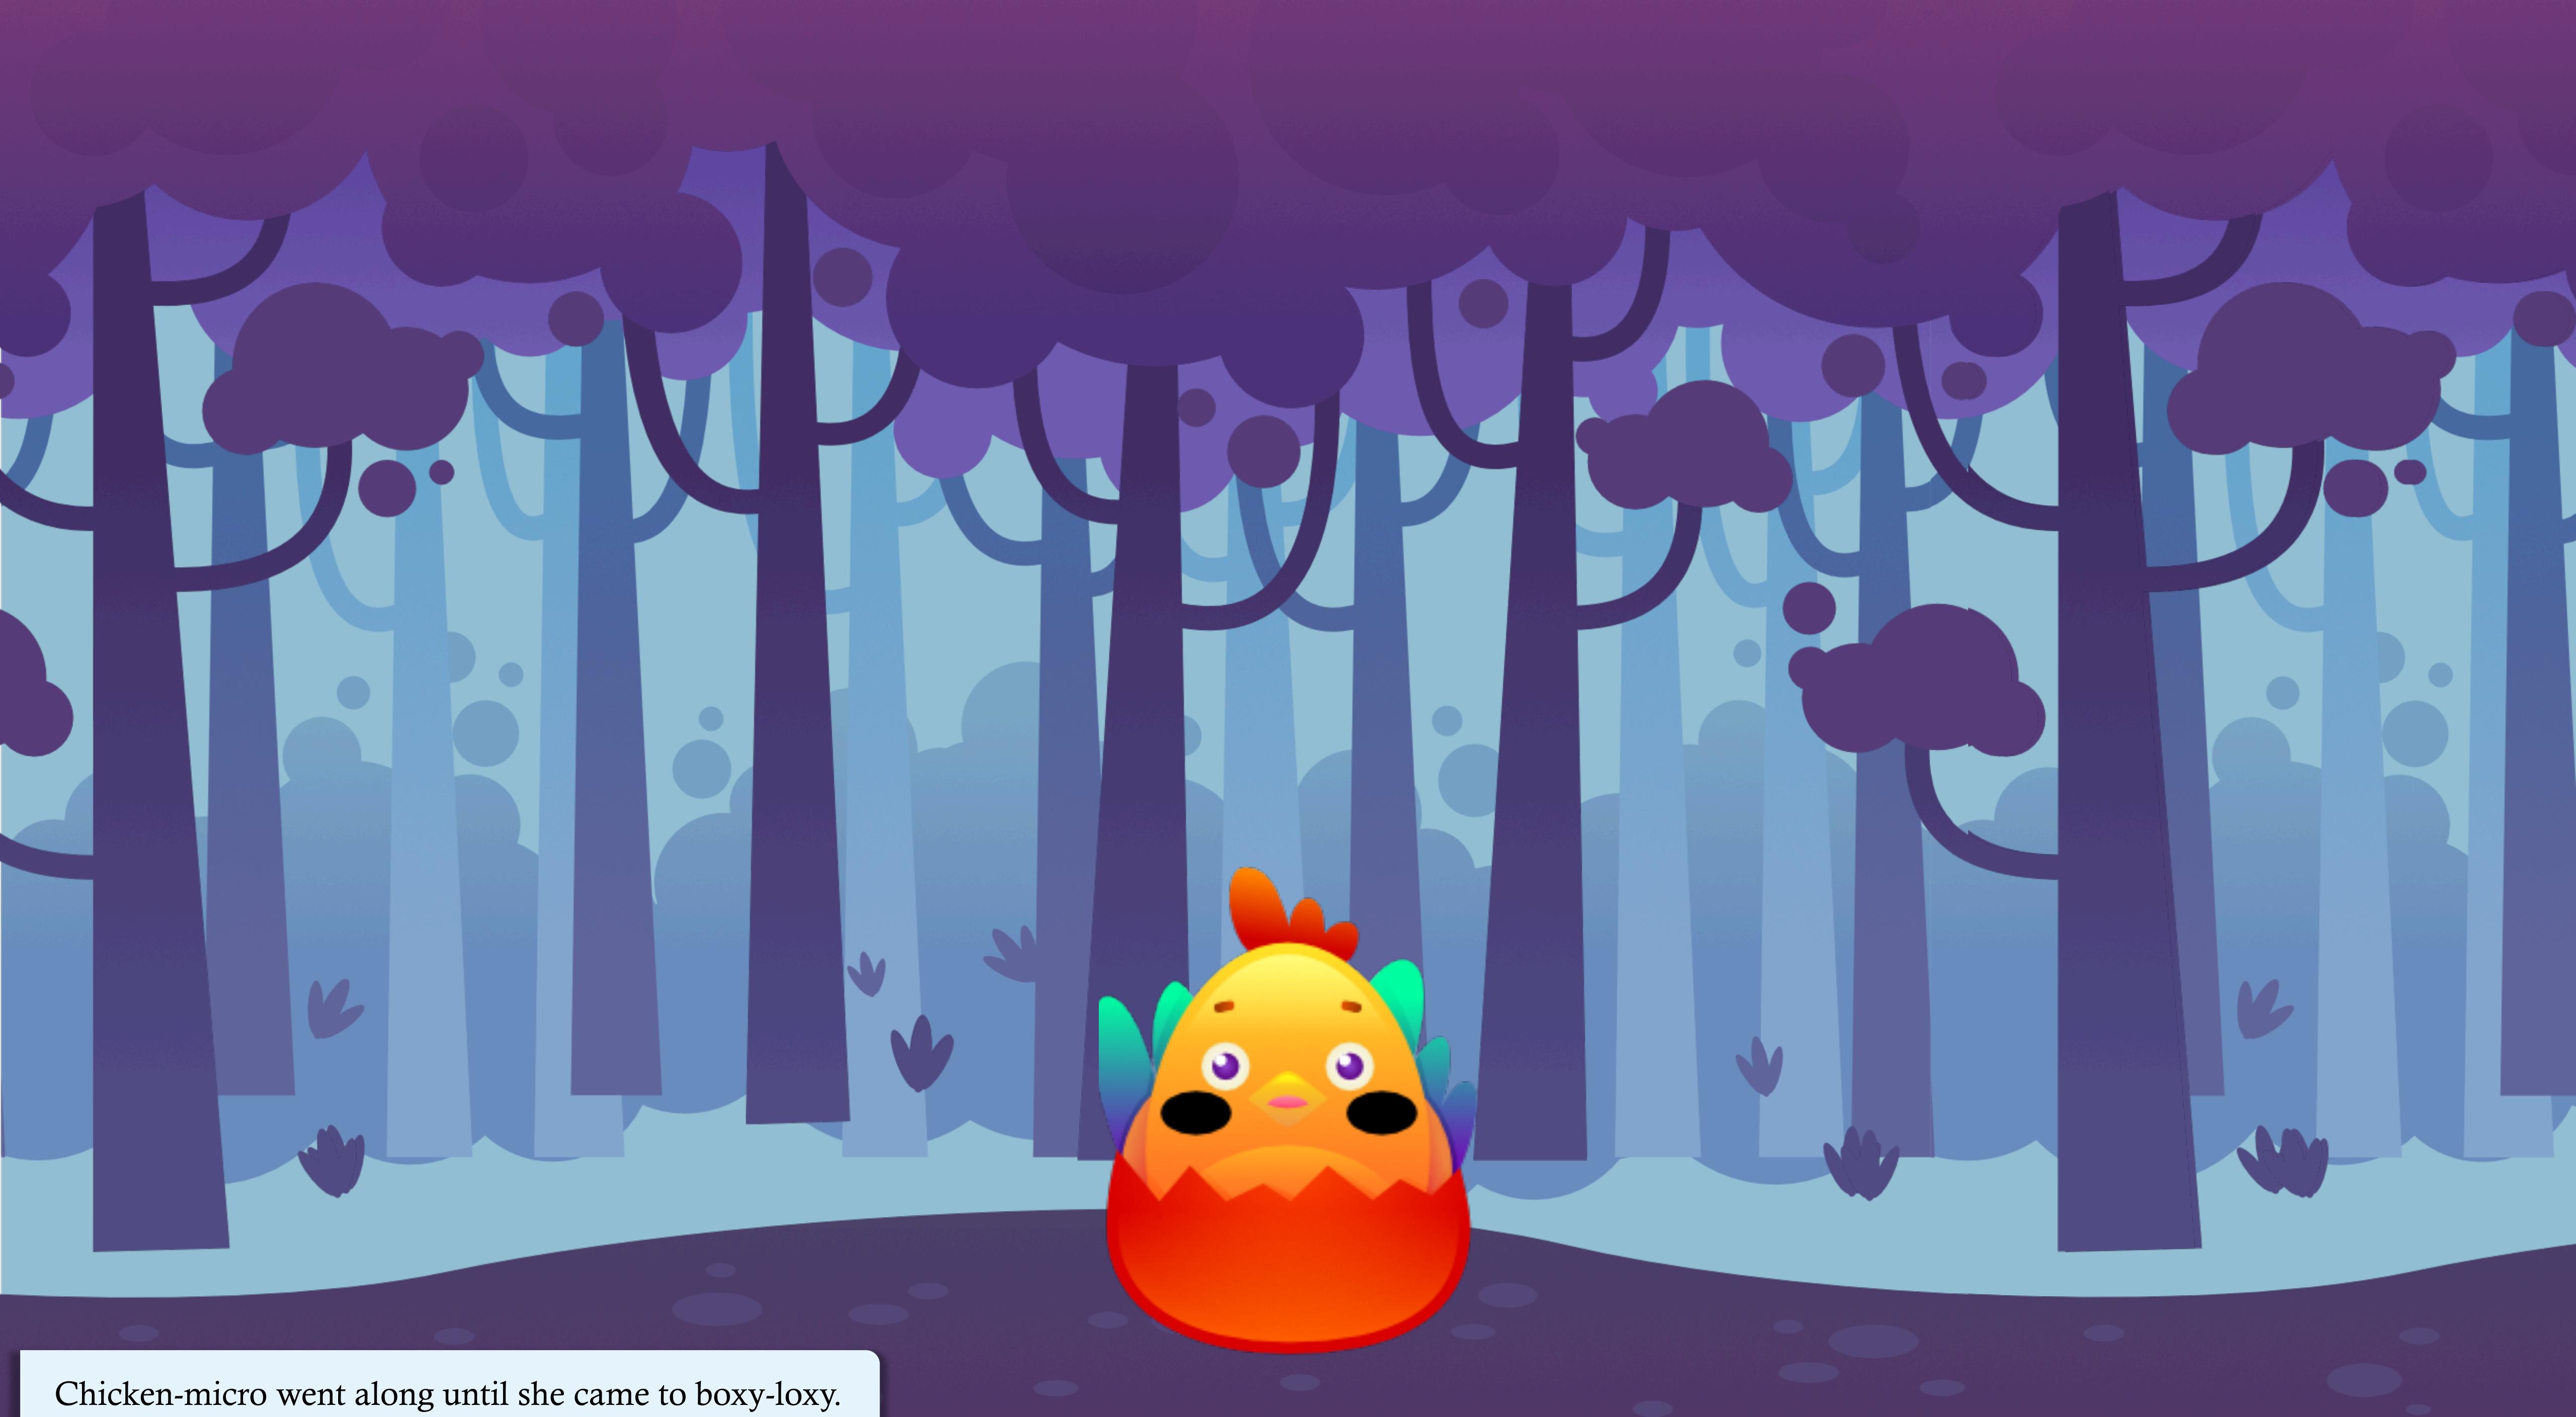
Wyatt Preul – jsgeek.com/nodesummit

One day Chicken-micro was making a request to the corn server
when—boom! The request timed out.

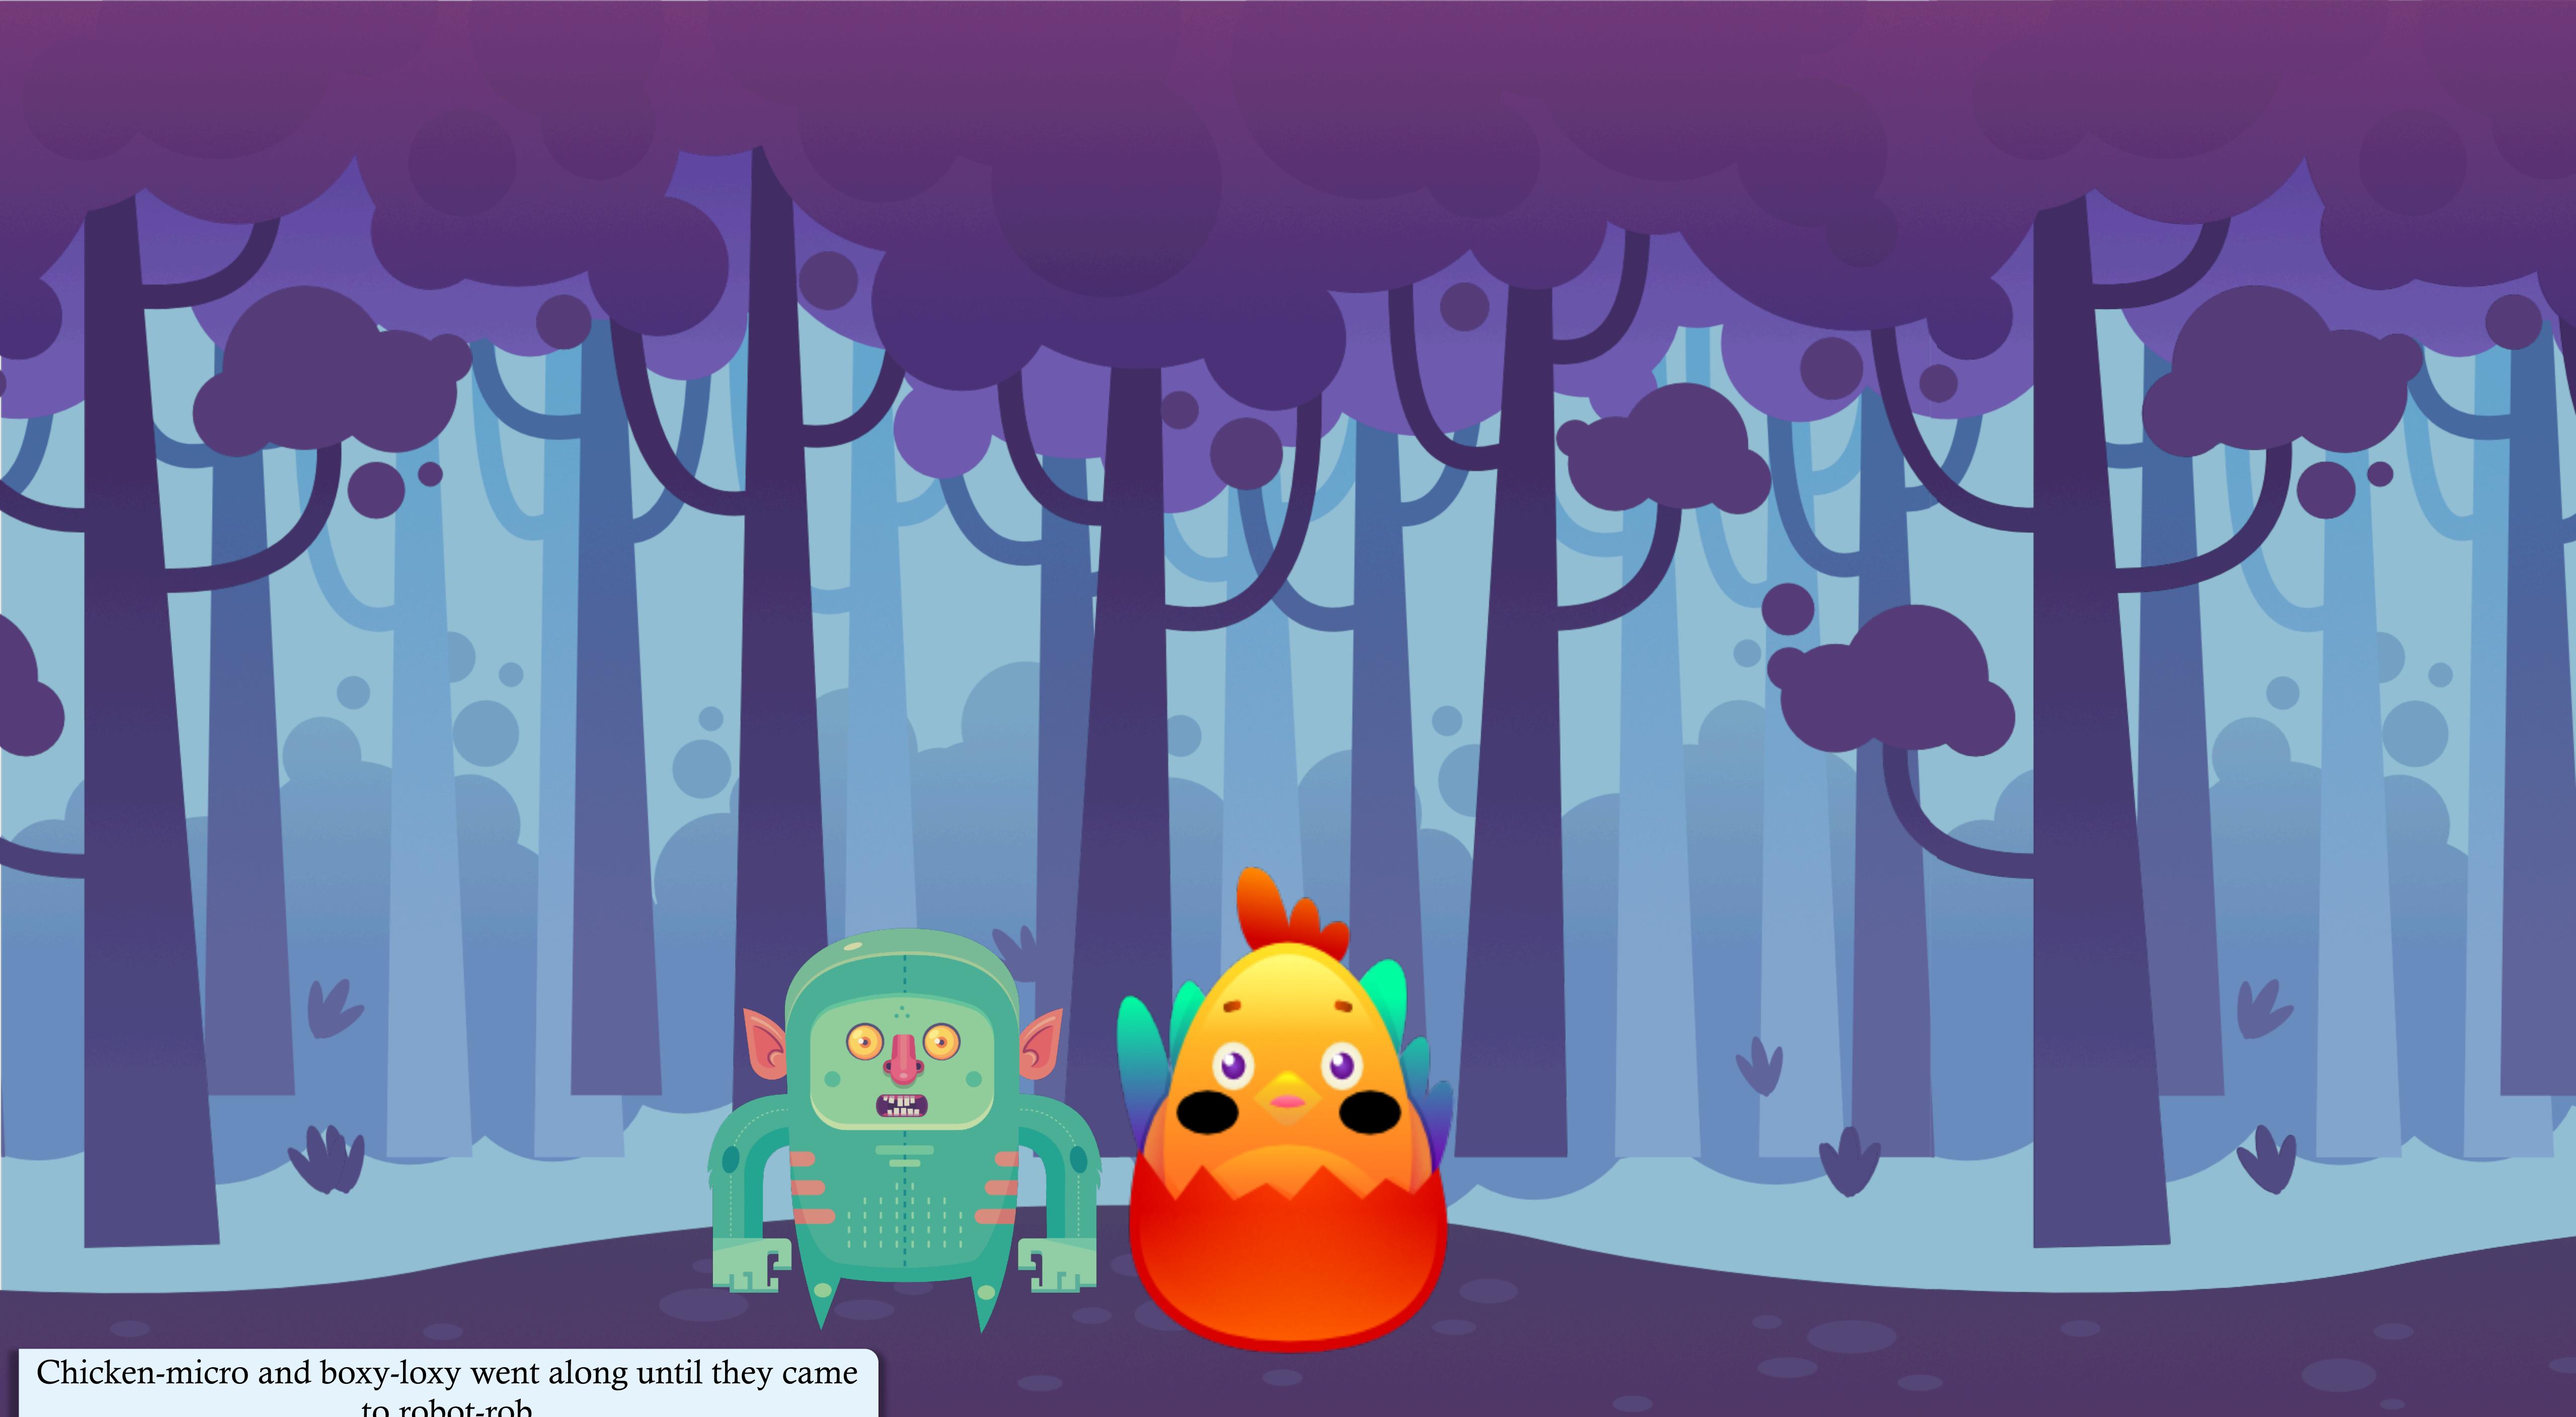




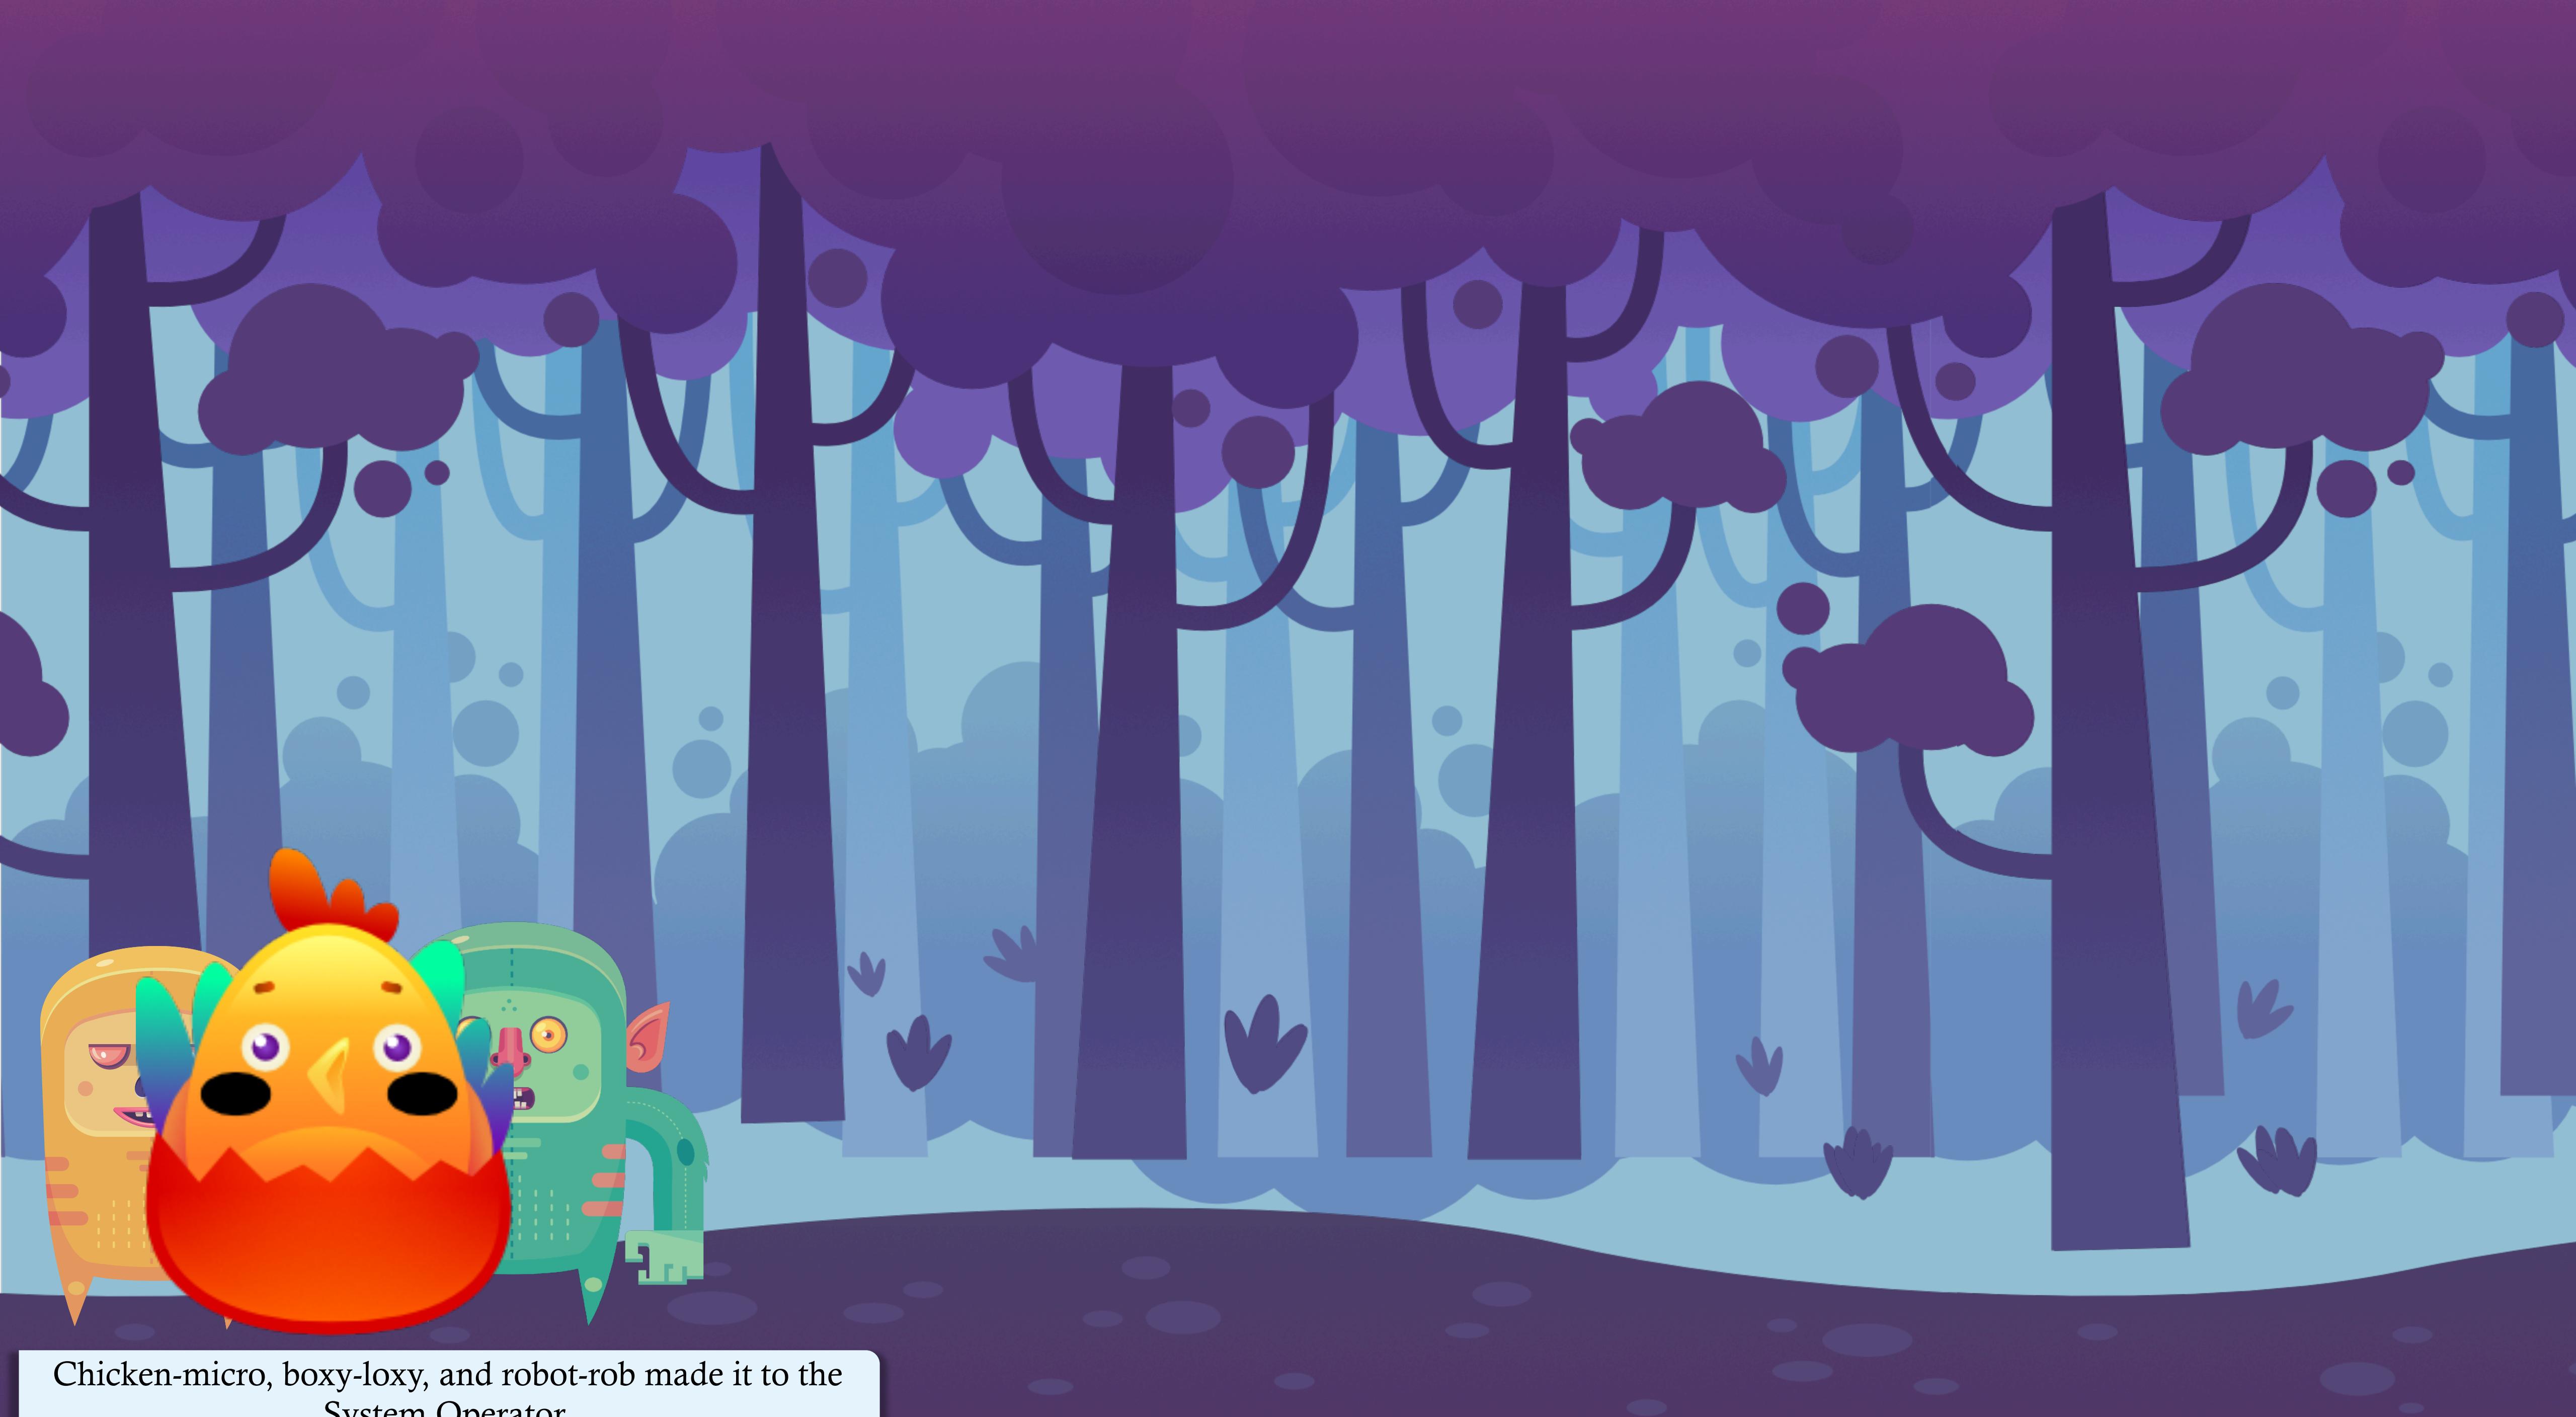
Goodness gracious, the servers
must be down. I must go and tell
the System Operator.



Chicken-micro went along until she came to boxy-loxy.



Chicken-micro and boxy-loxy went along until they came
to robot-rob



Chicken-micro, boxy-loxy, and robot-rob made it to the
System Operator.

Chicken-micro Postmortem

- Overloaded server should have returned 503, load balancer should have taken server out of rotation
- How would chicken-micro know that there are other healthy servers and that the next request will succeed?
- When microservices are architected the same as the services they replace, they will face the same challenges as the previous services.

What is a Microservice?

- Small, decomposed from monolith
- Isolated and independently deployable services
- Stateless and less fragile when changes are introduced

Benefits of Microservices

- Embrace failure, works in spite of external failures
- Iterate quickly - disposable services, independently deployable services

Microservice AntiPatterns

- Load balancer between microservices
- Startup order matters
- + many others, Microservices AntiPatterns and Pitfalls by Mark Richards

Load balancer between microservices

```
const getCorn = function (id, cb) {
  const url = `http://corn.prod.site.com/${id}`;
  wreck.get(url, { timeout: 5000 }, (err, res, corn) => {
    if (err && err.message === 'Client request timeout') {
      // retry the request
      return setTimeout(() => getCorn(id, cb), 1000);
    }
    // ... handle err or success
  });
};
```



Load balancer between microservices

```
const getCorn = function (id, cb) {
  const url = `http://corn.prod.site.com/${id}`;
  wreck.get(url, (err, res, corn) => {
    // server is under load, retry
    if (err && err.output.statusCode === 503) {
      return setTimeout(() => getCorn(id, cb), 1000);
    }
    // ... handle err or success
  });
};
```



Startup order matters

```
const mysql = require('mysql');

const connection = mysql.createConnection({
  host: process.env.DB_HOST
});

connection.connect();
// no fallback for db not found
```





Autopilot Pattern

- Apps that can be deployed and scaled with a single click.
- Apps and workflows that work the same on our laptops as in the cloud (public and private cloud).
- Apps and workflows that aren't married to any specific infrastructure or scheduler.
- Further reading at autopilotpattern.io

Autopilot Applications

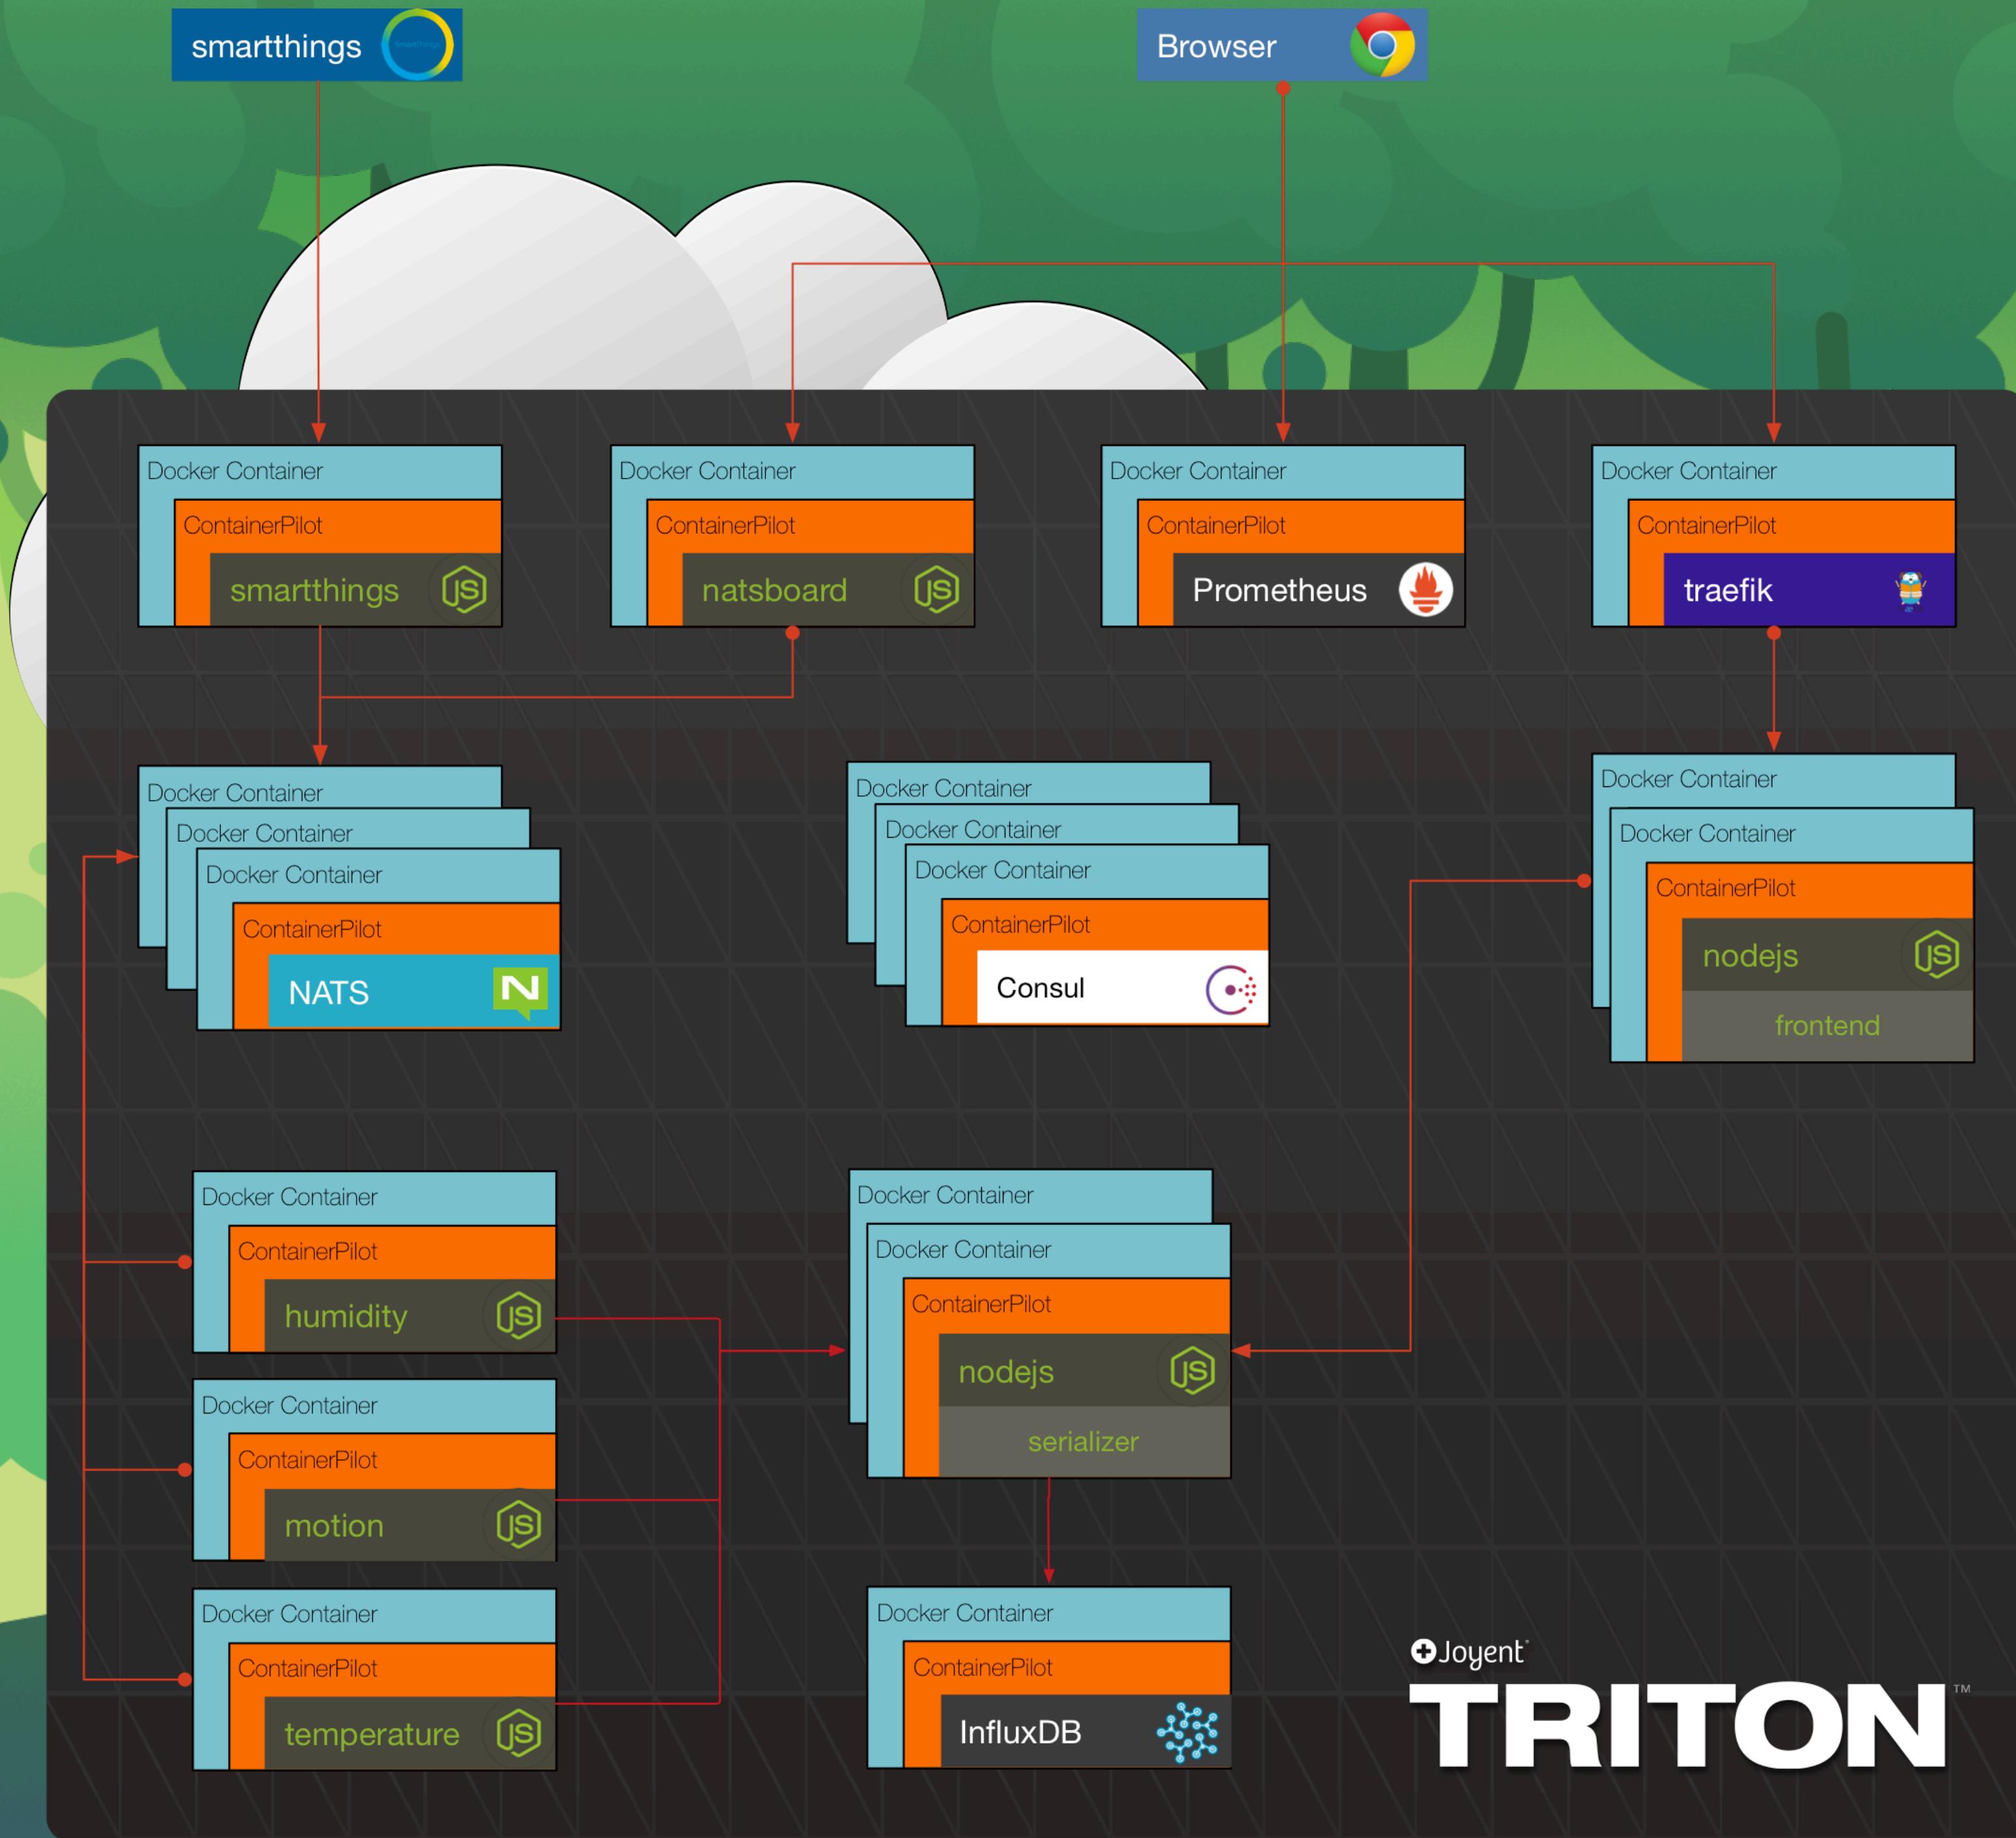
- github.com/autopilotpattern - solutions that follow the Autopilot Pattern
- MongoDB, MySQL, InfluxDB, Consul, NATS, Wordpress, Jenkins, ...
- Container support through ContainerPilot

Autopilot in Practice

- Applications composed of portable docker containers
- Service discovery through consul or another catalog
- Container local health and services respond to service dependency changes

Node.js Example

- <https://github.com/autopilotpattern/nodejs-example>
- Modules used:
 - hapi
 - Seneca
 - Piloted
 - Wreck





Isolation

```
function setupDb () {
  const influxServer = Piloted.service('influxdb');
  if (!influxServer) {
    internals.db = bufferedDb;
    return;
  }

  internals.db = new Influx.InfluxDB({
    host: influxServer.address,
    port: influxServer.port
  });
}
```

Change

```
Piloted.on('refresh', () => {
  setupDb();
});
```

ContainerPilot

- Automates service discovery, life cycle management, and telemetry reporting
- Capabilities
 - Container-local health checks
 - PID 1 init process
 - Service discovery registration and watchers
 - Telemetry reporting to Prometheus
 - Free & Open Source! github.com/joyent/containerpilot

ContainerPilot.json5 Watches

```
watches: [
  {
    name: 'influxdb',
    interval: 3
  }
],
jobs: [
  {
    name: 'onchange-influxdb',
    exec: 'pkill -SIGHUP node',
    when: {
      source: 'watch.influxdb',
      each: 'changed'
    }
  }
]
```

ContainerPilot.json5 Jobs

```
{  
  name: 'serializer',  
  port: {{.PORT}},  
  exec: 'node /opt/app/',  
  health: {  
    exec: `/usr/bin/curl -o /dev/null --fail  
      -s http://localhost:{{.PORT}}/health`,  
    interval: 2,  
    ttl: 5  
  }  
}
```

hapi Circuit Breaker

```
const server = new Hapi.Server({
  load: {
    sampleInterval: 50          // milliseconds
  }
});

server.connection({
  port: process.env.PORT,
  load: {
    maxEventLoopDelay: 20      // milliseconds
  }
});
```

Client-side Load Balancing

```
// Round-robin
const service = Piloted.service('nats');

// You decide
const services = Piloted.serviceHosts('nats');

// Grab a random healthy service
services[Math.floor(Math.random() * services.length)];
```



Load Balancers at the Edge

- Don't expose microservices directly outside of your organization
- Setup a load balancer that can use Consul, in our example this is traefik, could be HA Proxy, nginx, or Express-Gateway
- API gateways have a place when exposing microservices externally

Telemetry Reporting

```
telemetry: {  
  port: 9090,  
  tags: ['op'],  
  metrics: [  
    {  
      namespace: 'example',  
      subsystem: 'process',  
      name: 'event_delay',  
      help: 'Node.js event loop delay',  
      type: 'gauge'  
    }  
  ]
```

hapi Metrics with Topsy



Further Reading

- All links can be found at: jsgeek.com/nodesummit
- Autopilot Pattern - autopilotpattern.io
- hapi - hapijs.com
- ContainerPilot – joyent.com
- Seneca – senecajs.org
- Tao of Microservices - Richard Rogers, Manning Press